

Appendix C: Interview Notes (Anonymised – Internal UBS)

Excerpts from Stakeholder Conversations

Overview

This appendix presents anonymised excerpts from the stakeholder conversations conducted for this inquiry. Three stakeholders participated: a technology peer (Stakeholder A), a team lead with management responsibilities (Stakeholder B), and a tech risk and compliance specialist (Stakeholder C). The conversations were semi-structured, exploring perceptions of AI use in communication, concerns about authenticity and trust, and expectations for governance.

Direct quotes are presented in italics. All identifying information has been removed to protect confidentiality.

Stakeholder A: Technology Peer

Context: Stakeholder A works in a technology role within the division. They use AI tools occasionally for documentation and drafting support.

On authenticity

When asked how they would feel receiving a message they knew was AI-generated, Stakeholder A said: "*It doesn't feel genuine. It doesn't feel representative of the person.*" They elaborated that AI-generated text often has a particular style. It tends to be verbose and overly polished. This is different from how most people actually write. "*You can usually tell. The language feels off.*"

On acceptance thresholds

Stakeholder A distinguished between partial and substantial AI involvement. "*Partially is okay, but if mostly it is worrying.*" They explained that using AI to fix grammar or restructure a paragraph is fine. But if AI writes most of the content, that changes things. The message stops being from the person and starts being from the machine.

On human oversight

Stakeholder A emphasised the importance of human review. "*The human needs to continuously refine and legitimise it.*" They were clear that AI output cannot be sent without careful checking. The leader remains responsible for accuracy and tone.

On AI accuracy

Stakeholder A expressed scepticism about AI reliability. They noted that AI can produce confident but incorrect statements. "*You can't just trust it. It makes things up.*" This concern informed the policy's emphasis on fact verification.

Stakeholder B: Team Lead

Context: Stakeholder B is a team lead with management responsibilities. They use AI tools regularly for meeting summaries, grammar checking, and drafting support.

On current AI use

Stakeholder B described AI as "*a useful thinking partner*" that helps "*20% of the time*." They use it primarily for low-risk tasks: summarising meetings, checking grammar, improving clarity. They noted that AI helps them sound more professional, particularly for complex or sensitive messages.

On disclosure

When asked whether leaders should disclose AI use, Stakeholder B gave a nuanced answer. "*They should but wouldn't.*" They explained that disclosure feels awkward. People might be embarrassed to admit they used AI or worried about how others would react. "*There's a stigma. People might think you can't write properly.*"

On cognitive dependency

Stakeholder B raised concerns about over-reliance on AI. "*My concern is using it for thinking for us.*" They worried that leaders might gradually stop developing their own arguments. AI would do the intellectual work, and leaders would just review and send. This could erode the skills that make leadership communication effective.

On governance expectations

Stakeholder B expected the organisation to provide clear guidance. They mentioned approved tools, data protection rules, and escalation channels. "*We need to know what's allowed. And what happens if something goes wrong.*" They also expressed interest in peer review mechanisms for high-stakes communications.

On the artefact (feedback stage)

When shown the draft policy, Stakeholder B responded positively to the three-tier structure. "*This makes it clear. You know where you stand.*" They appreciated the decision framework as a practical tool. However, they raised concerns about enforcement. "*Some people will just ignore it if there's no consequence.*" They also questioned the absolute prohibition on AI for HR matters, suggesting that limited AI assistance for framing developmental conversations might be helpful.

Stakeholder C: Tech Risk and Compliance Specialist

Context: Stakeholder C works in a tech risk and compliance role with management responsibilities. They bring a governance perspective to AI use in communication.

On regulatory considerations

Stakeholder C emphasised the compliance dimension. In financial services, communication errors can have regulatory consequences. "*If something goes out with wrong information, we need to know who's accountable. AI doesn't change that.*" They stressed that regulatory bodies expect human accountability for all client and market-facing communications.

On data protection

Stakeholder C was particularly concerned about data entering external AI systems and UBS internal AI assistants. "*Client data cannot go into ChatGPT or any other AI tool used in UBS. That's non-negotiable.*" They noted that many employees may not realise the data protection implications of pasting information into AI tools. Clear boundaries are essential.

On approved tools

Stakeholder C expected a clear list of approved AI tools. "*We need to know which tools have been vetted. Which ones meet our security requirements.*" They noted that enterprise versions of AI tools often have different data handling arrangements than consumer versions. The policy needs to distinguish between them.

On escalation

Stakeholder C emphasised the need for clear escalation pathways. "*People need to know who to ask when they're not sure. Otherwise, they'll guess, and sometimes they'll guess wrong.*" They suggested that Compliance should be the escalation point for any communication with regulatory implications.

Summary of Key Themes

Across the three conversations, several themes emerged consistently. All stakeholders accepted AI as a support tool but drew a line when AI does most of the work. All expressed concerns about authenticity and the risk that AI-generated communication feels impersonal. All expected human oversight to be meaningful, not a formality. The disclosure dilemma was acknowledged by all: disclosure is appropriate but culturally awkward. And all expected governance infrastructure: approved tools, data protection, escalation, and clear accountability.

These themes directly informed the AI Communication Policy. The artefact was designed to address the gaps and concerns stakeholders identified, translating their insights into practical guidance that could work in the real organisational context.