

Appendix B: Stakeholder-Policy Mapping

How Inquiry Findings Shaped the AI Communication Policy

Overview

This appendix documents the connection between stakeholder insights and policy design. Each section shows how a specific finding from the inquiry informed one or more elements of the AI Communication Policy. The goal is to demonstrate that the artefact is grounded in real organisational evidence, not abstract principles.

1. Conditional Acceptance of AI

What stakeholders said

Stakeholder A stated: "Partially is okay but mostly is worrying." The stakeholders expressed comfort with AI as a support tool but discomfort when AI does most of the work. There is an implicit threshold. Cross it, and trust erodes.

How this shaped the policy

This finding directly informed the three-tier classification system. Permitted uses cover situations where AI polishes human work. Restricted uses cover situations where AI contributes substantively but disclosure is required. Prohibited uses cover situations where AI involvement is inappropriate regardless of disclosure. The structure acknowledges that the acceptability of AI depends on degree, not just presence.

2. Authenticity Concerns

What stakeholders said

Stakeholder A noted that AI-generated content "doesn't feel genuine" and "doesn't feel representative of the person." The stakeholders observed that AI tends to produce verbose, overly polished language that differs from natural communication styles. Recipients may question whether the leader actually stands behind the words.

How this shaped the policy

This finding shaped the Authenticity principle, which states that some communications require genuine human voice. It also informed the verification requirements, which include checking tone and removing AI artefacts like generic phrases and filler sentences. The prohibited uses section identifies contexts where authenticity is non-negotiable, such as performance feedback and conflict resolution.

3. The Disclosure Dilemma

What stakeholders said

Stakeholder B said leaders "should but wouldn't" disclose AI use. The stakeholder's believed disclosure was appropriate in certain situations but acknowledged that embarrassment, fear of judgement, and uncertainty would prevent it. There is a cultural discomfort with admitting AI involvement.

How this shaped the policy

This finding shaped the Transparency principle and the restricted uses category. By making disclosure a formal requirement in specific contexts, the policy normalises transparency. It removes individual stigma by establishing an organisational expectation. The suggested disclosure statement ("This document includes AI-assisted drafting and has been reviewed for accuracy") is deliberately neutral to reduce embarrassment.

4. Human Oversight as Non-Negotiable

What stakeholders said

Stakeholder A emphasised that "the human needs to continuously refine and legitimise" AI output. The stakeholders insisted that AI cannot replace judgement. Review must be meaningful, not a formality.

How this shaped the policy

This finding shaped the Accountability principle, which states that leaders own everything sent under their name regardless of AI involvement. It also informed the detailed verification requirements. The five-step checklist ensures that review is substantive: read the full message, verify facts, check tone, confirm technical accuracy, and remove AI artefacts. This addresses the concern that human review might become a rubber stamp.

5. Cognitive Dependency Risk

What stakeholders said

Stakeholder B expressed concern about leaders "using it for thinking for us." The worry was that AI assistance might gradually replace human intellectual effort. Leaders might stop developing their own arguments and rely on AI to do the thinking.

How this shaped the policy

This finding reinforced the framing throughout the policy that AI is for augmentation, not replacement. The permitted uses section emphasises that AI should polish work, not create it. The decision framework includes the question "Would the recipient expect my authentic voice?" to prompt reflection on whether AI is doing too much. The verification requirements ensure that leaders engage critically with AI output rather than accepting it passively.

6. Governance Infrastructure Expectations

What stakeholders said

The stakeholders expected strong organisational controls. They mentioned approved AI tools, strict data protection boundaries, clear escalation channels, peer review, and transparency about model provenance. This reflects financial services culture where governance is expected.

How this shaped the policy

This finding informed several policy sections. The Data Protection principle prohibits entering confidential information into unapproved systems. The Definitions section specifies approved tools. The Roles and Responsibilities section assigns ownership to Technology Leadership for maintaining the approved tools list. The Escalation section provides clear pathways for uncertain situations. The quarterly review cycle ensures the policy evolves with changing technology and governance needs.

7. AI Error and Accuracy Concerns

What stakeholders said

Stakeholder A noted scepticism about AI accuracy, particularly its tendency to produce confident but incorrect content. There was awareness that AI "makes things up" and that leaders cannot simply trust AI output.

How this shaped the policy

This finding directly informed the verification requirements. The policy explicitly states that leaders must verify facts against real sources because AI can generate plausible but incorrect information. The decision framework includes the question "Could AI errors cause harm or reputational damage?" to prompt extra caution in high-stakes contexts. The language in the verification section is deliberately blunt: "AI makes things up."

Summary Table

The table below provides a condensed view of the mapping between stakeholder findings and policy sections.

Stakeholder Finding	Policy Section(s)	How It Was Addressed
Conditional acceptance ("partially ok, mostly worrying")	Sections 6.1, 6.2, 6.3	Three-tier classification system based on degree of AI involvement
Authenticity concerns ("doesn't feel genuine")	Sections 5.4, 6.3, 9	Authenticity principle, prohibited uses for sensitive contexts, tone verification
Disclosure dilemma ("should but wouldn't")	Sections 5.1, 6.2	Transparency principle, normalised disclosure requirements with neutral language
Human oversight non-negotiable	Sections 5.2, 9	Accountability principle, detailed verification checklist
Cognitive dependency risk	Sections 6.1, 8, 9	Augmentation framing, decision questions, active verification requirements
Governance infrastructure expectations	Sections 4, 5.3, 7, 10, 12	Approved tools, data protection, roles, escalation, quarterly review
AI error and accuracy concerns	Sections 8, 9	Decision question on harm, mandatory fact verification